

CS 6476 Project 3

Sai Sandeep Dasari

sdasari38

903540744





Gradescope Group Quiz Collaboration Photo

<Insert a picture of your group here>





The screenshot shows a web browser window displaying the Gradescope interface for a 'Project 3 - Group Quiz'. The browser's address bar shows the URL: `gradescope.com/courses/155064/assignments/753851/submissions/new`. The page header indicates '0/5 Questions Answered'. The main content area is titled 'Project 3 - Group Quiz' and displays 'Q1' worth '1 Point'. The question text is: 'Select the incorrect option about HarrisNet from the following:'. There are four radio button options:
1. 'The second-order gradient matrix can be calculated from the first-order gradient matrix through simple matrix multiplication operations.'
2. 'Whether or not a pixel is a corner is decided based on the corner response function.'
3. 'Sometimes in a window there might be many valid corner points detected; since these points all pass the strict detection criteria, we should keep them all in consideration.'
4. 'Image gradient layers can be calculated through simple convolutions based on Sobel filters.'
A 'Save Answer' button is located below the options. A notification bar at the bottom of the browser window states 'bluejeans.com is sharing your screen.' with 'Stop sharing' and 'Hide' buttons. The bottom of the image shows a video call interface with three participants: 'Rishabh Dattani', 'Ashwin Lele', and another participant. The Windows taskbar is visible at the very bottom.

basic concepts of the

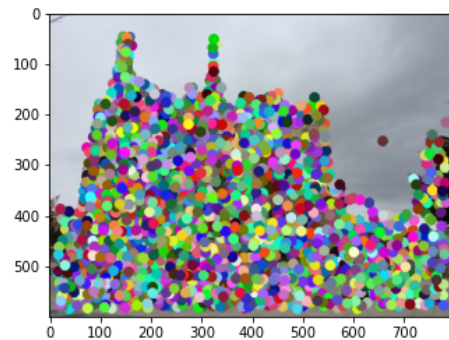



Part 1: HarrisNet



1000 (above) vs 4500 (below)



Part 1: HarrisNet



Part 1: HarrisNet

<Describe how the HarrisNet you implemented mirrors the original harris corner detector process. (First describe Harris) What does each layer do? How are the operations we perform equivalent?)>

Harris corner detection is an algorithm for detecting corners in an image which can be later used for a wide variety of tasks invariant to rotation. The HarrisNet we implemented mirrors the original harris corner largely in terms of what the layers of our network do:

1. ImageGradientsLayer : calculate image gradients in each direction (X and Y), just like the original implementation
2. ChannelProductLayer: In Szeliski 4.1.1. Algorithm 4.1, this step computes the other products of the above gradients. We compute product between channels of the previous layer: I_{xx} , I_{yy} and I_{xy} .
3. SecondMomentMatrixLayer: compute Second Moment Matrix. We convolve the above 3 images (3 channels) with a Gaussian too
4. CornerResponseLayer: computes the R cornerness matrix, the R cornerness matrix is the equivalent “scalar interest metric” in the original
5. NMSLayer : performs non-maxima suppression to keep only the strongest corners in local regions. Finally the original algorithm simply finds local maxima above a certain threshold. We do the same with our layer.

Part 2: SiftNet

<Describe how the SiftNet you implemented mirrors the Sift Process. (First describe Sift) What does each layer do? How are the operations we perform equivalent?>

Scalar Invariant Feature Transform (SIFT) is a feature descriptor for calculating low dimensional vectors in an image that are scale invariant. In our implementation

We do not perform the sub-octave difference in Gaussian layers in our implementation. Instead we find Image Gradients and extract gradient information along each orientation direction (Lowe used 36bins, in our case 8).

Then we use the orientations to create weighted histograms over the entire image. This step mirrors the original process. Once we find the $8 \times 16 = 128$ SIFT feature vector, we create feature vectors that accumulate histograms from local regions.

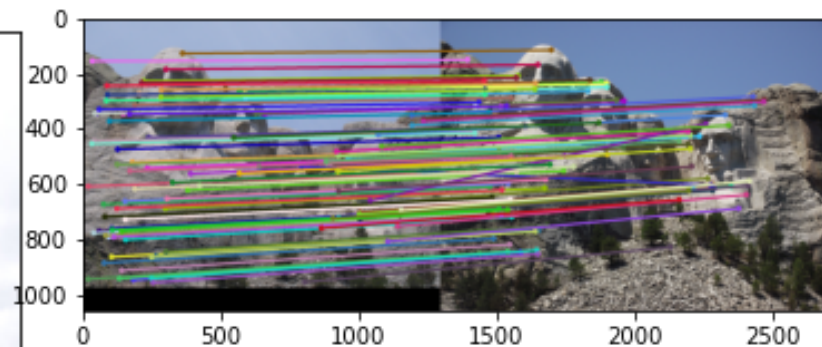

Part 2: SiftNet

- <Explain what we would have to do make our version of Sift rotationally invariant (conceptually)>
- <Explain what we would have to do to make our version of SIFT scale invariant (conceptually)>


Part 2: SiftNet

- <What would happen if instead of using 16 subgrids, we only used 4 (dividing the window into 4 grids total for our descriptor)?>
- This will make it difficult for us to detect local features effectively as the size of the window is greatly reduced. There will be mismatches (compared to groundtruth) in the final matching as the small local regions in various portions of the image are matched (all kinds of circular wheels are easily matched)
- <What could we do to make our histograms in this project more descriptive?>
- We can add more orientations in the histogram that will improve the feature vector by adding a few more dimensions.

Part 3: Feature Matching



Part 3: Feature Matching





<Describe your implementation of feature matching.>


In feature matching, we initially used our keypoints detected from the HarrisNet and discarded some points.

We did this by finding a distance vector between the 2 images and using the ratio test with a threshold of 0.8. Thus, a good number of features were found at the same rejecting mismatches minimally. The result was
Notredame – 0.92
Rushmore – 0.95 in accuracy.

Results: Ground Truth Comparison



Results: Ground Truth Comparison



<Insert numerical performances on each image pair here.>

Notredame : 100/100 required matches
Accuracy = 0.92

Rushmore: 100/100 required matches
Accuracy = 0.95

Gaudi: 71/100 required matches
Accuracy = 0.000000

Results: Discussion

- <Discuss the results. Why is the performance on some of the image pairs much better than the others?>
- The model did well on Notredame and Rushmore because the 2 images A and B have almost the same amount of lighting where as Gaudi's A and B differed largely different in terms of brightness and contrast. Also the Gaudi image B is a smaller resolution image thus the model struggles to find more interest points below the ratio threshold.
- <What sort of things could be done to improve performance on the Gaudi image pair?>
- Firstly the resolution of both images needs to be brought to a uniform amount before running Harris Corner. Also, some pre-processing to bring the color brightness and contrast as close as possible will help define the texture of the surfaces and consequently find strong matches.

Conclusions

<Describe what you have learned in this project. Feel free to include any challenges you ran into.>

This project is a great intro into feature descriptors and matching. This project also allows to clearly understand the progression between Harris corners and SIFT descriptors. The project also helped establish the use of neural networks by using PyTorch's neural network. It was particularly difficult to debug the values flowing through the network even though the weights are defined by us.

The application of deep learning to solve this problem is quite clear and evident from this project.

Extra Credit: Sift Parameter variations

Extra Credit: Custom Image Pairs